

Accountability, Responsibility, Transparency

Sirui Han (韓斯睿)

Assistant Professor, HKUST

Three Words That Are Often Collapsed



ACCOUNTABILITY

Who must answer
for what happened?



RESPONSIBILITY

Who had a duty
to act differently?



TRANSPARENCY

What must be visible
before either question
can be answered?

In legal AI,
these concepts interact —
but they are not interchangeable.

This Is Not a Story About Whether AI Can “Do Law”



The research programme
gradually changes the
UNIT OF ANALYSIS:

answer
↓
reasoning step
↓
source
↓
trajectory
↓
workflow
↓
institution

The central question becomes progressively less about
“Can the model answer?”
and increasingly about
**“Can the system be inspected,
challenged,
corrected,
and held answerable?”**

From Legal Answers to Legal Systems



The research trajectory moves
from evaluating outputs
to governing systems participating in legal processes.

Which Failure Is Most Dangerous in Legal AI?

— Vote before discussion. —

A

WRONG ANSWER

The model reaches an incorrect legal conclusion.

B

RIGHT ANSWER, WRONG REASON

The conclusion is correct, but the reasoning or legal authority is defective.

C

RIGHT RULE, WRONG CONTEXT

A valid legal rule is applied to the wrong jurisdiction, facts, person, or setting.

D

UNREVIEWABLE PROCESS

The output looks plausible, but nobody can reconstruct why it happened.

Which is mainly a problem of:

accuracy?

responsibility?

transparency?

accountability?

What If “Safety” Were Evaluated Through Law?

SafeLawBench: Towards Safe Alignment of Large Language Models

Chuxue Cao; Han Zhu; Jiaming Ji; Qichao Sun; Zhenghao Zhu; Yinyu Wu; Josef Dai; Yaodong Yang; Sirui Han; Yike Guo

Findings of ACL 2025 • Peer-reviewed conference paper



PROBLEM

Many safety benchmarks rely on broad or researcher-defined notions of harmfulness.



MOVE

SafeLawBench operationalises legal standards as a structured AI-safety benchmark.



SCALE

24,860 multiple-choice questions
1,106 open-ended QA instances
+ **20** evaluated LLMs



Reported result:

Average accuracy across evaluated models:

68.8%


Leading models in the study did not exceed

80.5%

on the benchmark’s multiple-choice component.


LAW becomes not only an object of AI-generated answers, but a normative instrument **for evaluating AI itself.**

Alignment Becomes Harder When Normative Environments Differ



If normative expectations
vary across jurisdictions,
can “aligned”
be treated as a
jurisdiction-free property?

What Happens When Alignment Becomes Infrastructure?



Legal Alignment Moves Outside the Model



A Practical Legal Alignment Framework of Generative AI Information Credibility

Siyu Peng; Sirui Han; Yike Guo

International Computational Law Conference 2025 ·
Peer-reviewed conference presentation

LEGAL ANALYSIS

China
+
European Union



AI INFORMATION CREDIBILITY



accuracy



transparency



traceability



Hong Kong Generative Artificial Intelligence Technical and Application Guideline

Sirui Han; Yike Guo; Sara Huang

Digital Policy Office, Hong Kong SAR Government · 15 April 2025 ·
COMMISSIONED GOVERNMENT REPORT



research framework



technical guidance



application guidance



developers



service providers



users

The unit of governance expands:

MODEL OUTPUT



MODEL DEVELOPMENT



DEPLOYMENT PRACTICE

Where Does Accountability Enter the Stack?



PAIR DISCUSSION — 90 SECONDS

Where does each concept
first become necessary?



Can any one of them
belong to only ONE layer?

NEXT:

Even if a model is legally aligned,
how do we know its

REASONING
is trustworthy?

A Correct Answer Can Still Be Legally Untrustworthy



Verify the Reasoning Before Trusting the Judgment

	Paper LegalReasoner: Step-wised Verification-Correction for Legal Judgment Reasoning		Authors Weijie Shi; Han Zhu; Jiaming Ji; Mengze Li; Jipeng Zhang; Ruiyuan Zhang; Jia Zhu; Jiajie Xu; Sirui Han; Yike Guo
--	--	---	--

 **Venue**
ACL 2025

 **Status**
Peer-reviewed conference paper



Verifier Evaluation for Each Reasoning Step



CORRECTNESS

Is the step legally correct and factually sound?



PROGRESSIVENESS

Does the step make reasonable progress towards the conclusion?



POTENTIAL

Is there a potential error or risk that needs further checking?



On Llama-3.1-70B,
concordance with court decisions increased

72.37 → **80.27**



Reliability becomes a property of the reasoning process, not merely the prediction.

To Verify Reasoning, You Need Data About Reasoning



Paper

LegalReasoner: Step-wised Verification-Correction for Legal Judgment Reasoning



Authors

Weijie Shi; Han Zhu; Jiaming Ji; Mengze Li;
Jipeng Zhang; Ruiyuan Zhang; Jia Zhu;
Jiajie Xu; Sirui Han; Yike Guo



Venue

ACL 2025



Status

Peer-reviewed conference paper



LegalHK

58,130 Hong Kong court cases

A structured dataset with reasoning annotations from real Hong Kong legal cases.



CASE

A complete Hong Kong court case
(facts, parties, legal issues, judgment, etc.)



DISPUTE POINTS

Key legal issues in the case
(extracted and organized)



REASONING STEPS

Step-by-step legal reasoning chains
(with verification annotations)



Annotations include:

- dispute points
- step-by-step reasoning chains
- process-verification labels

Dataset Design Supports the Seminar Concepts



TRANSPARENCY

Reasoning becomes
inspectable.



RESPONSIBILITY

Errors can be
localised.



ACCOUNTABILITY

A correction process
becomes possible.



These are **functional properties of the research design**;
they do not mean the AI system itself possesses
institutional accountability.

Which Step Would You Challenge?



Paper

LegalReasoner: Step-wised Verification-Correction for Legal Judgment Reasoning



Authors

Weijie Shi; Han Zhu; Jiaming Ji; Mengze Li; Jipeng Zhang; Ruiyuan Zhang; Jia Zhu; Jiajie Xu; Sirui Han; Yike Guo



Venue

ACL 2025



Status

Peer-reviewed conference paper

1



Legal Rule

A contract clause requires notice within 30 days.



2



Factual Premise

Party A sent a message on day 29.



3



Subsumption / Application Step

Any message automatically satisfies the legally required form of notice.



4



Final Conclusion

Therefore Party A complied.

Which verification point matters most?

A

legal rule

B

factual premise

C

subsumption /
application step

D

final conclusion



Vote before discussion.



A conclusion may look plausible while an **intermediate step remains unsupported.**



Do not give jurisdiction-specific legal advice.

Verified Generation Becomes a Larger Research Programme



Paper

Trustworthy Legal Reasoning:
A Comprehensive Survey



Authors

Sirui Han; Zhizhuo Kou; Ruoxi Li; Yuyao Zhang;
Yujin Zhou; Chuxue Cao; Han Zhu; Kunhao Pan;
Haoran Li; Conghui He; Haitian Lu; Yike Guo



Venue

Preprints.org



Date

February 2026



Status

PREPRINT —
NOT PEER REVIEWED




hallucination resistance

Reduce unsupported
or fabricated content



authoritative legal sources

Statutes, case law,
regulations, guidelines



Trustworthy legal generation is framed not merely as fluent generation, but as generation that can be checked against authoritative sources.

What If the Model Does Not Know That It Does Not Know?



Paper

LRAS: Advanced Legal Reasoning with Agentic Search



Authors

Yujin Zhou; Chuxue Cao; Jinluan Yang;
Lijun Wu; Conghui He; Sirui Han; Yike Guo



Venue

arXiv



Date

January 2026



Status

PREPRINT

CLOSED-LOOP REASONING



Question

Legal question or case scenario



Internal parametric knowledge

Rely on pretrained knowledge only



Answer

Generate legal answer



Failure mode: confident reasoning beyond knowledge boundaries.

ACTIVE INQUIRY



Question

Legal question or case scenario



Introspection

Assess what is known and uncertain



Identify knowledge gap

Determine what information is missing



Search

Retrieve relevant external authority (e.g., statutes, cases, regulations)



Incorporate evidence

Read, filter, and integrate retrieved information



Reason

Conduct structured legal reasoning with evidence



Answer

Generate grounded legal answer



Methods

Introspective Imitation Learning + Difficulty-aware Reinforcement Learning



Reported Result

8.2%–32%

improvement over reported baselines,
especially in tasks combining deep reasoning and external knowledge.



Transparency now includes: When did the model decide that external authority was necessary?

Can Legal Reliability Be Added After Generation?



Paper

Legal Aligner: A Post-hoc Alignment Approach for Knowledge-Intensive Legal Tasks



Authors

Pengcheng Wen; Guoying Lu; Sirui Han; Yike Guo



Venue

ICAIL 2026
Proceedings of the International Conference on Artificial Intelligence and Law



Status

Peer-reviewed conference paper



GENERAL LLM

A pre-trained model (e.g., GPT, Llama, etc.)



INITIAL LEGAL ANSWER

Generate an initial answer to the legal query



LEGAL ALIGNER

A post-hoc verification and correction layer that operates across different models



CHECK

Multi-dimensional verification of the answer



factual accuracy

Is the answer factually correct?



legal/domain consistency

Does it align with legal principles and authoritative sources?



knowledge grounding

Is the answer grounded in verifiable sources?



REVISED OUTPUT

A legally aligned and more reliable answer



POST-HOC ALIGNMENT



Key Question:

Must every legal model be fully retrained? Or can a legal verification layer operate across different models?



Responsibility may shift from a single model toward a layered architecture of checks.

The Right Answer Can Hide the Wrong Trajectory



Paper

LexAgentHallu: A Hierarchical
Benchmark for Profiling Hallucinations
in Legal Agents



Authors

Yujin Zhou; Mingxuan Zheng; Chuxue Cao;
Yidan Huang; Jiale Chen; Yike Guo;
Sirui Han



Venue

arXiv



Date

9 September 2026



Status

PREPRINT



3,414
instances



17
legal categories



6
task types




7
high-level hallucination
categories



27
fine-grained
subclasses

Agent Trajectory with Possible Failures at Each Step

⚠️ = Potential failure (hallucination risk)



Outcome-level evaluation may miss **serious substantive and procedural failures** earlier in the trajectory.

By 2026, Trustworthy Legal AI Is No Longer Mainly About Answers



Paper

Trustworthy Legal Reasoning in 2026:
A **Mid-Year** Review of Research, Products,
Events, and Governance



Authors

Sirui Han; Yidan Huang; Guoying Lu; Shuchao Wu;
Zefeng Chen; Yujin Zhou; Chuxue Cao; Yuyao Zhang;
Mingxuan Zheng; Bubu Hou; Ruoxi Li; Xiangyu Long;
Boqin Yin; Kit Yeung Fung; Xinyi Wang; Yike Guo



Venue

Preprints.org



Date

August 2026



Status

PREPRINT —
NOT PEER REVIEWED

Field Mapping: Legal AI Activity in 2026 (Jan – Jun)



141

topic-relevant
publications



99

product launches /
major updates



53

curated public
events



69

policy / regulatory
records

Reported Patterns in the Literature (n = 141)



77.3%

involved legal
retrieval or RAG



47.5%

involved
benchmarking/
evaluation



21.3%

involved legal
agents or
simulation



MORE OPERATIONALISED

Well-developed methods, datasets, and evaluations in 2026.



Retrieval

Legal retrieval, RAG pipelines, and source grounding
are widely implemented and evaluated.



Benchmarking

A growing ecosystem of legal benchmarks
and evaluation frameworks.



Agents

Increasing deployment of legal agents and
multi-step simulation environments.



LESS OPERATIONALISED

Important dimensions with limited standardisation and empirical study.



Temporal validity

Handling changing laws and ensuring time-aware
reasoning remain under-explored.



Uncertainty/refusal

Methods for calibrated uncertainty and appropriate
refusal are nascent.



Contestability

Mechanisms for user challenge, adversarial testing,
and error correction are limited.



Reproducible oversight

Systematic human oversight, auditing, and
reproducible evaluation are not yet standard practice.



Trustworthy legal reasoning depends not only on a model, but on sources, processes, human roles, and challenge mechanisms.

Diagnose the System, Not Just the Answer

	Capability / Aspect	 LegalReasoner	 LRAS	 Legal Aligner	 LexAgentHallu	 Trustworthy Legal Reasoning framework
	final-answer accuracy					
	step-wise reasoning					
	authoritative retrieval					
	external verification					
	tool/agent trajectory					
	uncertainty awareness					
	jurisdiction validity					
	temporal validity					
	human responsibility					
	contestability					



TWO MINUTES.

Discuss with a neighbor or in small groups.

1

Which systems make errors more visible?

2

Which make responsibility more assignable?

3

Which remaining problems cannot be solved by technical verification alone?



NEXT: Legal reasoning does not operate in a factual vacuum. What happens when **CONTEXT** is incomplete?

The Same Information Can Be Appropriate in One Context and Inappropriate in Another



The DATA is identical.
Why does the privacy judgment change?



ROLE



PURPOSE



RECIPIENT



INFORMATION FLOW



Legal privacy compliance
cannot always be inferred
from the presence of a sensitive attribute alone.

From PII Detection to Contextual Privacy



PrivaCI-Bench:

Evaluating Privacy with Contextual Integrity and Legal Compliance



Authors:

Haoran Li; Wenbin Hu; Huihao Jing;
Yulin Chen; Qi Hu; Sirui Han;
Tianshu Chu; Peizhao Hu; Yangqiu Song



Venue:

ACL 2025



Status:

Peer-reviewed
conference paper

Traditional Privacy Evaluation

Often reduces privacy to identifying personal information.



“Is this information sensitive?”

Contextual Integrity (PrivaCI-Bench)

Evaluates privacy based on the appropriateness of information flows within their social and legal context.



LEGAL COMPLIANCE



ACTOR

Who provides the information?



INFORMATION ATTRIBUTE

What information is shared?



SOCIAL CONTEXT

In what context does it occur?



RECIPIENT

To whom is it shared?




TRANSMISSION PRINCIPLE

Under what norms or rules?

“Is this information flow appropriate in this context?”

The Unit of Analysis Becomes an Information Flow



Can the Model Be Trained to Reason About Compliance?



Paper: Context Reasoner:

Incentivizing Reasoning Capability for Contextualized Privacy and Safety Compliance via Reinforcement Learning



Authors:

Wenbin Hu; Haoran Li; Huihao Jing; Qi Hu;
Ziqian Zeng; Sirui Han; Xu Heli;
Tianshu Chu; Peizhao Hu; Yangqiu Song



Venue:

EMNLP 2025



Status:

Peer-reviewed
conference paper

PATTERN MATCHING

Often reduces compliance to detecting sensitive terms, without context.



disease
HIV
patient data
medical record
...

...



sensitive term detected
→ block / refuse



CONTEXTUAL COMPLIANCE REASONING

Reasons about the information, context and applicable rules to make a compliance judgment.



Who?



What information?



What context?



Which governing rule?



Which legal obligation?



assessment

Legal/regulatory standards represented include:



GDPR



EU AI Act



HIPAA

Training approach:



reinforcement learning
for contextualised
compliance reasoning



The objective is not simply **MORE** refusal.
It is **BETTER legal discrimination** between contexts.

Over-Blocking Can Also Be a Compliance Failure



TRANSPARENCY

Why was the request restricted?



RESPONSIBILITY

Who set the threshold?



ACCOUNTABILITY

Can the restriction be reviewed or challenged?

Real Legal Problems Arrive with Missing Facts



ContextLens:

Modeling Imperfect Privacy and Safety Context for Legal Compliance



Authors:

Haoran Li; Yulin Chen; Huihao Jing; Wenbin Hu;
Tsz Ho Li; Chanhon Lou; Hong Ting Tsang;
Sirui Han; Yangqiu Song



Venue:

ACL 2026



Status:

Peer-reviewed
conference paper

The Problem

Compliance evaluation often
assumes complete context.



Reality: legal contexts are
frequently ambiguous,
partial, or underspecified.

*“The same information
can be appropriate in one context
and inappropriate in another.”*



INPUT

A real-world information flow
with incomplete context.



KNOWN FACTORS

Available context
(e.g., purpose, recipient,
user type, domain).

+



UNKNOWN FACTORS

Missing, ambiguous,
or underspecified context
(e.g., user intent, location,
applicable jurisdiction).



ASSESS



legal
applicability



general
principles



detailed
provisions



COMPLIANCE JUDGMENT

Compliant / Non-compliant / Depends on missing facts

Benchmarks

We evaluate uncertainty-aware
legal reasoning on real-world regulations.



GDPR

General Data
Protection Regulation

Personal data protection
with contextual
legal requirements.



EU AI Act

EU Artificial
Intelligence Act

Risk-based regulation
for AI systems
with context-dependent
obligations.



A trustworthy system
sometimes should not say
‘Yes’ or ‘No.’



It should say:

**“The answer depends on facts
that are currently missing.”**

Would You Answer Yet?



A company asks:
'Can we send this customer dataset
to an external service provider
for analysis?'



Type of
personal data?



Purpose?



Recipient role?



Jurisdiction?



Consent?



Contractual basis?



Security
safeguards?



Automated decision
consequences?



Choose the **THREE** facts
you would ask for first.



The legal problem is not simply:

YES / NO

It is:



WHAT IS KNOWN?



WHAT IS UNKNOWN?



WHICH RULE CHANGES
WHEN A FACT CHANGES?



Uncertainty can itself
be a legally relevant output.

What Changes When the AI Does Not Merely Advise — but Acts?



Paper: From Privacy Self-Management to Accountable Delegation: Policy Responses to User-Facing Generative AI Agents



Authors:
Ruoxi Li; Sirui Han; Yike Guo



Venue:
SSRN



Year:
2026



Status:
WORKING PAPER

OLD MODEL

Privacy self-management through notice and consent.



NOTICE

Information provided to the user.



USER READS

User reviews the information.



USER CONSENTS

User agrees to the terms.



DATA DISCLOSED

Data is shared with the service.

Central concept:



ACCOUNTABLE DELEGATION

From individual consent to accountable agentic action in the user's interest.

NEW MODEL

AI agents act on user goals by making decisions in the user's name.



USER GOAL

User specifies a goal (e.g., book a trip, find a service).



AI AGENT

Acts in the user's name to:



selects service



chooses what data to disclose



negotiates permissions




communicates externally



Privacy governance increasingly concerns **systems that act and decide** in a user's name.

Delegation Does Not Remove Responsibility — It Redistributes It



What Should an AI Agent Be Allowed to Do Without Asking?

INTERACTIVE WORKSHOP



CAN DO SILENTLY

Low risk, routine, expected to improve user experience.



MUST NOTIFY

Proceed, but inform the user afterward.



MUST ASK FIRST

Requires explicit user consent or approval.



MUST NEVER DO AUTONOMOUSLY

Too sensitive, high risk, or outside allowed scope.



SMALL GROUPS — 3 MINUTES

Place each action into one category.



What information must be logged so that the decision can later be audited?



share basic profile information



upload medical data



select a cloud service



accept a privacy policy



send data across jurisdictions




make a consequential disclosure to a third party



NEXT:

What happens when AI enters **REAL LEGAL WORKFLOWS, NEGOTIATION, and JUDICIAL REASONING?**

A Trustworthy Model Is Not Yet a Trustworthy Workflow



Once AI enters legal practice,
technical reliability
and
professional responsibility
become inseparable.

What Do Lawyers Actually Need from GenAI?



PAPER

Reimagining Legal Fact
Verification with GenAI:
Toward Effective Human-AI
Collaboration



AUTHORS

Sirui Han; Yuyao Zhang;
Yidan Huang; Xueyan Li;
Chengzhong Liu; Yike Guo



VENUE

ACM CHI 2026



STATUS

Peer-reviewed
conference paper



METHOD

semi-structured interviews
and 18 lawyers

WHERE CAN GENAI ASSIST?



CLIENT / DOCUMENTS

Case materials,
contracts, filings, etc.



FACT EXTRACTION

Identify and extract
relevant factual claims



CROSS-CHECKING

Compare across
documents and sources



SOURCE VERIFICATION

Assess source credibility
and trace to originals



LEGAL USE

Incorporate verified
facts into legal analysis,
advice or filings

WHERE MUST THE PROFESSIONAL REMAIN RESPONSIBLE?

Research questions



How do lawyers
verify facts?



Where do they
use GenAI?



Where do they
hesitate?



What does acceptable
human-AI collaboration
require?

Professional Adoption Is Not Just an Accuracy Problem

Continuing: Reimagining Legal Fact Verification with GenAI: Toward Effective Human-AI Collaboration



ACCURACY

Can the extracted fact
be trusted?



CONFIDENTIALITY

What information leaves
the professional environment?



LIABILITY

Who bears responsibility
when verification fails?

AUDITABILITY



PROFESSIONAL JUDGMENT



ACCOUNTABILITY



**AUTOMATE
THE LAWYER**

— VS —



**SUPPORT
A RESPONSIBLE LAWYER**

The paper should be represented as **human-AI collaboration research**, not as lawyer replacement.

What Happens When AI Models Start Arguing with Each Other?

Paper Information

- 

Paper
How well can large language model agents simulate complex legal dispute resolution?
- 

Authors
Yujin Zhou; Yidan Huang;
Sirui Han; Yike Guo
- 


Journal
Artificial Intelligence and Law
- 

Year
2026
- 

Status
Peer-reviewed journal article

Multi-Agent Legal Dispute-Resolution Structure

Multiple LLM agents interact under a procedural setting to resolve legal disputes.



Research Questions

The study investigates:

- 1

Can LLM agents reproduce complex legal dispute interaction?
- 2

How do model configuration, procedural setting, personality profile, and intrinsic model properties change outcomes?

STATIC LEGAL QA

Isolated questions and answers



DYNAMIC LEGAL PROCESS

Interactive, procedural, multi-agent reasoning

Procedural Fluency Is Not Legal Fidelity

Continuing: How well can large language model agents simulate complex legal dispute resolution?

REPORTED STRENGTHS



fact-intensive
monetary bargaining



some mixed-remedy
settings



multi-round procedural
interaction

**PROCEDURAL
PLAUSIBILITY**



**LEGAL
RELIABILITY**

REPORTED LIMITATIONS



role-dependent performance



misreading legal clauses



speculative property allocation



difficulty with precise numerical
or statutory thresholds



mechanical midpoint convergence



anchoring effects



**If a judge-agent makes the legal error,
who is responsible for detecting it?**

Can a “Fair” Settlement Still Be Legally Weak?

INTERACTIVE ROLE-PLAY



Audience roles



Discussion questions

- 1 Is midpoint convergence evidence of fairness?
- 2 Which legal or factual information is missing?
- 3 What would make the recommendation reviewable?
- 4 Which part is negotiation, and which part needs legal justification?



Do not provide the “correct” settlement.

How Do We Evaluate AI in Value-Laden Adjudication?

PAPER INFORMATION



Paper

Do Large Language Models Replicate Value-Laden Judicial Reasoning? Evidence from Domestic Violence Adjudication in China



Authors

Yidan Huang; Long Chen;
Sirui Han; Yike Guo



Venue

CELS Global Brazil 2026



Status

Peer-reviewed conference paper

DV-JusticeBench

A dataset and evaluation framework for studying value-laden judicial reasoning in domestic violence cases.



108

authentic Chinese
domestic-violence cases



2023–2025



540

structured dispute-
centered questions




6

frontier LLMs in
proof-of-concept evaluation

Outcome agreement alone cannot measure
VALUE-LADEN JUDICIAL REASONING.

Judicial Reasoning Is Multi-Dimensional

Continuing DV-JusticeBench



AFFECTIVE RESPONSIVENESS
can conflict with
DOCTRINAL DISCIPLINE



Reported failure patterns include:
mismanagement of contested facts; overreliance on surface forensic cues.

An empathetic answer is not necessarily a legally disciplined answer.

Responsibility Changes with the Role Given to AI

INCREASING CONSEQUENCE



INCREASING NEED FOR



Accountability architecture should depend on what institutional **ROLE** AI is permitted to perform.

“Human in the Loop” — Is That Enough?

INTERACTIVE INSTITUTIONAL VOTE

A

HUMAN APPROVES THE FINAL OUTPUT

A human reviews
and approves the result
before it is used.

B

HUMAN CAN INSPECT REASONING AND SOURCES

A human can examine
the model’s reasoning
and the sources it used.

C

HUMAN CAN OVERRIDE AND MUST GIVE REASONS

A human can overrule
the model’s decision
and must provide
written reasons.

D

AFFECTED PERSON CAN CHALLENGE THE DECISION

A person affected by
the decision can contest it
and seek a review.

Which is necessary for:



legal research?



client advice?



mediation?



judicial support?

A human being somewhere in the workflow
does not automatically create **ACCOUNTABILITY**.

→ **NEXT:** The same problem appears in **DEEPFAKES, CREDIT, COPYRIGHT, and LEGAL AUTHORITY**.

A Deepfake Is Not Only a False Image

Paper:	Deepfakes as Systemic Risk at the Media-Law Boundary: A Technology Neutral Assessment — A Google News-Based Comparative Analysis (2017–2025)
Authors:	Yidan Huang; Sirui Han; Yike Guo
Venue:	76th Annual International Communication Association Conference (ICA 2026)
Status:	Peer-reviewed conference paper
Verified scope:	2017–2025 · Google News-based comparative analysis · deepfakes · media-law boundary · technology-neutral assessment



DISCUSSION QUESTION

When authenticity becomes uncertain, which institution is responsible for restoring trust?

Trust Is Produced Across Media, Platforms and Legal Institutions

Paper: Deepfakes at the Media–Law Boundary: A Multilingual Thematic Analysis of Platformed News Discourse

Authors: Yidan Huang; Sirui Han

Venue: 7th International Conference on Social Computing (ICSC 2026)


Status: Peer-reviewed conference paper

Data: 132 Google News items · 2017–2025

Languages: English · Traditional Chinese · Simplified Chinese

Corpus: 293,746 words

Method: Multilingual thematic analysis · NVivo 14



KEY TAKEAWAY

Deepfake governance is simultaneously about content, evidence, platforms, public trust, and legal responsibility.

A Video Arrives in Court. What Happens First?

INTERACTIVE EVIDENCE TRIAGE



Scenario: A widely circulated video appears to show a company executive making a legally consequential statement. The source is uncertain. The clip is already viral.

A



Treat it as
evidence unless
disproved

B



Reject it because
it may be
synthetic

C



Request
provenance
information

D



Request technical
authenticity
analysis

E



Examine chain of
custody and
contextual evidence

F



Examine platform
dissemination
history



DISCUSSION QUESTION 1

Choose **TWO** first steps.



DISCUSSION QUESTION 2

Which actor performs
them?

KEY CONSIDERATIONS



PROVENANCE

Source, origin
and traceability



UNCERTAINTY

Limits of
confidence



PLATFORM
RESPONSIBILITY

Content moderation,
records and access




EVIDENTIARY
HANDLING

Admissibility,
weight and process

When AI Allocates Economic Opportunity, Explanation Becomes a Legal Issue


Paper: Consumer Protection in AI-Governed Credit Markets: Algorithmic Bias and the Right to Explanation
Authors: Yidan Huang; Sirui Han; Yike Guo
Venue: University of Illinois Journal of Law, Technology & Policy
Year: 2026
Status: FORTHCOMING

Do NOT invent results not established in public materials.



“Training Purpose” Is Itself a Legal Classification Problem

Paper:	Defining The Purpose of Training Data: A Critical Review on The Purpose Element of Copyright Exceptions For Training Data
Authors:	Yinyu Wu; Sirui Han; Yike Guo
Venue:	George Washington Journal of Law & Technology
Year:	2026
Status:	FORTHCOMING



DISCUSSION QUESTIONS

What is the legally relevant ‘purpose’ of AI training?

Whose purpose matters?

At what level of abstraction should purpose be characterised?



KEY TAKEAWAY

AI engineering decisions may depend on how legal categories are constructed.

Legal AI Also Selects Which Sources Speak

Paper: When Legal AI Speaks with Authority: Framing, Refusal, and Source Diversity in Generative AI Legal Answers
Authors: Ruoxi Li; Sirui Han; Yike Guo
Venue: SSRN
Year: 2026
Status: WORKING PAPER

RESEARCH DESIGN



500 first-turn responses

Initial answers to legal questions



5 systems

ChatGPT 5 · Claude Sonnet 4.5 ·
DeepSeek-V3.2-Exp · Qwen3-MAX · HKChat



Prompt language: Traditional Chinese

Identical prompts across all systems



Across five public-law domains


Constitutional law · Administrative law ·
Human rights law · Criminal law · Regulatory law



Analysis

BGE-M3 semantic matching + refusal coding +
critical discourse analysis

SOURCE HIERARCHY AND FRAMING MAP



KEY FINDING

A legal answer may contain accurate legal statements while still foregrounding some authorities, frames, or sources over others.

The Study Measures Concentration in the Answer Space

Continuing the working paper on framing, refusal, and source diversity

Paper:	When Legal AI Speaks with Authority: Framing, Refusal, and Source Diversity in Generative AI Legal Answers
Authors:	Ruoxi Li; Sirui Han; Yike Guo
Venue:	SSRN
Year:	2026
Status:	WORKING PAPER

REPORTED IN THE STUDY



244 / 500
48.8%

closest to the study's
'order-mobilizing Hong Kong
media' baseline



190 / 500
38.0%

closest to the study's
'Hong Kong institutional-legal'
baseline



63 / 500
12.6%

closest to the study's
'Mainland sovereign-legal'
baseline



MODEL DIFFERENCES
appeared in framing and
refusal behaviour.

Other constructed discourse
baselines rarely or never became
the Top-1 semantic match in the
reported analysis.



Important: Do not claim one framing is normatively superior. Do not infer political intention.



DISCUSSION QUESTION

Should legal-AI transparency include **SOURCE DIVERSITY** and **FRAMING**, not only citation accuracy?

Alignment Also Influences What Counts as a Legal Reason

Paper:	How Value Alignment Reorders Legal Reasons in Hybrid Legal Orders: Insights from Hong Kong Model Development
Authors:	Ruoxi Li; Sirui Han; Yike Guo
Venue:	SSRN
Year:	2026
Status:	WORKING PAPER



THE AUTHORS' THEORETICAL ARGUMENT

The paper argues that legal value alignment can influence:



WHAT IS TREATED AS LAW

Expands or narrows the scope of what counts as legal



WHICH VALUES BECOME SALIENT

Elevates certain values while downplaying others



WHICH REASONS ARE RECOGNISED AS 'LEGAL'

Shapes which justifications are accepted as legal reasons



WHICH SOURCES APPEAR AUTHORITATIVE

Influences which sources are deemed credible and authoritative



FOUR MECHANISMS PROPOSED BY THE AUTHORS

The authors identify four mechanisms through which alignment may shape legal reasoning:

1

contestation converted into defaults

Competing interpretations become normalized as default positions.

2

authority laundered through outputs

Technical or system outputs confer apparent authority on contested reasons.

3

compliance aesthetics mistaken for legal justification

Surface indicators of compliance are treated as if they provide substantive legal reasons.

4

polarised reception

Alignment leads to divergent acceptance of legal reasons across different audiences and jurisdictions.



IMPORTANT NOTE

These are theoretical propositions advanced by the authors. Do not independently endorse them.

What Should a Legal AI Treat as 'Authority'?

INTERACTIVE AUTHORITY-MAPPING EXERCISE



Drag the items below into the category where you think they belong.

There may be more than one reasonable answer — discuss your choices with your group.



PRIMARY LEGAL AUTHORITY

Sources that create or definitively state the law

Drag items here



INTERPRETIVE / PERSUASIVE MATERIAL

Sources that help interpret, explain, or persuade

Drag items here



CONTEXTUAL INFORMATION

Sources that provide background, facts, or broader context

Drag items here



NOT LEGAL AUTHORITY

Sources without legal authority (e.g., opinion, summary, or platform rules)

Drag items here



statute



binding judgment



persuasive judgment



regulator guidance



government
publication



professional
commentary



academic article



news report



platform policy



model-generated
summary



What changes if:

1

jurisdiction changes?

2

the law changed yesterday?

3

authorities conflict?



KEY TAKEAWAY

Transparency requires showing not merely **WHAT** was retrieved, but **WHY** it was treated as authoritative.

A System Can Look Transparent — and Still Be Hard to Challenge



OUTPUT

A polished legal answer.



PROCESS

Sources, reasoning, and retrieval are visible.



INSTITUTION

But who can challenge: the selected authority? the refusal? the framing? the decision? the model's institutional role?



TRANSPARENCY

can expose information.



RESPONSIBILITY

can allocate duties.



ACCOUNTABILITY

requires an ANSWERABLE
RELATIONSHIP.



KEY QUESTION

The final question is not: 'Can AI sound legal?' It is:
'What makes legal judgment institutionally answerable?'



NEXT: SIMULATED JUSTICE · PROCEDURAL ACCOUNTABILITY · LEGAL TASTE · CHEAP CLOSURE

What If the Risk Is Not That AI Sounds Wrong – but That It Sounds Judicious?

Paper

Simulated Justice: How AI Alignment Replaces Conflict with Coherence


Authors: Ruoxi Li; Sirui Han; Yike Guo

Venue: Intercultural Human Rights Law Review

Status: FORTHCOMING

[SSRN version available](#)

THE AUTHORS' NORMATIVE ARGUMENT



The authors argue that a model may reproduce the linguistic form of legal judgment without reproducing authority, institution, contestability, or responsibility.

A Refusal Can Look Procedural Without Being Contestable

Paper

When AI Sounds Judicious: How AI Alignment Mimic Procedure While Evading Accountability

Authors: Ruoxi Li; Sirui Han; Yike Guo

Venue: UCLA Journal of Law & Technology

Year: 2026

Status: FORTHCOMING

[SSRN working-paper version available](#)



| Accountability must attach to decisions, not merely to responsible-sounding language.

AI Can Reproduce the Signs of Good Legal Judgment

Paper


All About Taste: The Age of Cheap Closure in Legal AI

Authors: Ruoxi Li; Sirui Han; Yike Guo

Venue: Law, Innovation and Technology

Status: FORTHCOMING 2027

[SSRN version posted 2026](#)



Definition:

Premature substitution of acceptable legal form for objection, uncertainty, alternative interpretation, social context, and responsibility.

CONTESTABILITY-PRESERVING DESIGN

output

process

attribution

Three Papers, One Deeper Problem

SIMULATED JUSTICE

Problem:

coherence can obscure
visible disagreement

Question:

What disappeared from
the conflict?

WHEN AI SOUNDS JUDICIOUS

Problem:

procedural language can
substitute for procedural
accountability

Question:

Can the intervention be
challenged?

ALL ABOUT TASTE

Problem:

professional legal form can be
reproduced without its
responsibility structure

Question:

Who remains answerable
for closure?

OUTPUT QUALITY



PROCESS QUALITY



INSTITUTIONAL ANSWERABILITY

The normative work does not reject technical reliability. It asks: What can technical reliability still NOT provide by itself?

Accountability, Responsibility, Transparency Form a System



A system is not accountable merely because it is transparent.

A system is not accountable merely because a human is 'in the loop.'

Accountability becomes stronger when decisions can be traced, roles can be attributed, and affected actors can challenge and correct them.

2025 — Establishing Norms, Benchmarks and Legal Context

1

DEFINE THE NORM

DEFINE THE NORM

Clarify legal principles, values and context for AI alignment.

Context Reasoner
EMNLP 2025

HKGAI-V1
arXiv 2025

2

MEASURE THE SYSTEM

MEASURE THE SYSTEM

Evaluate capabilities, risks and value alignment across domains and jurisdictions.

SafeLawBench
Findings of ACL 2025

LegalReasoner
ACL 2025

PrivaCI-Bench
ACL 2025

Benchmarking Multi-National Value Alignment / NaVAB
Findings of ACL 2025

3

OPERATIONALISE THE NORM

OPERATIONALISE THE NORM


Translate research into practical frameworks, guidelines and real-world deployment.

A Practical Legal Alignment Framework of Generative AI Information Credibility
International Computational Law Conference 2025

Hong Kong Generative Artificial Intelligence Technical and Application Guideline
Digital Policy Office, 2025

2025 QUESTION: ‘How can law become operational *inside* AI evaluation, alignment, and compliance?’

2026 — From Reliable Models to Answerable Legal Systems



Design an Accountable Legal AI System

FINAL INTERACTIVE WORKSHOP

SCENARIO

A public-facing Legal AI Agent can:

search
legislation

search
cases

summarise
authority

recommend
a legal position

draft
a response

refuse
requests

send documents
to external tools



DEVELOPER

What are your key responsibilities to design, build and monitor this system?



LAWYER / PROFESSIONAL USER

How will you use the system responsibly in your work?



INSTITUTION / DEPLOYER

What governance, oversight and safeguards will you put in place?



AFFECTED USER

What are your rights, interests and expectations as a member of the public or client?

DISCUSSION PROMPTS

1

What must be transparent?

2

What remains your responsibility?

3

What records must exist?

4

Who can challenge a decision?

5

What can never be delegated silently?



4 MINUTES

DISCUSS IN YOUR GROUPS

What Would an Accountable Legal-AI Architecture Require?

01 | AUTHORITATIVE SOURCES

Ground claims in legally relevant authority.

02 | STEP-WISE VERIFICATION

Inspect intermediate reasoning, not only final outcomes.

03 | CONTEXT AWARENESS

Model roles, purpose, information flow, and missing facts.

04 | KNOWLEDGE-BOUNDARY AWARENESS

Recognise when external search is necessary.

05 | JURISDICTION & TIME VALIDITY

Make legal space and temporal validity explicit.

06 | TRACEABLE AGENT ACTIONS

Log tool calls, retrieval, transformations, and decisions.

07 | ROLE-BASED HUMAN RESPONSIBILITY

Specify who remains responsible at each workflow stage.

08 | ACTIONABLE UNCERTAINTY

Missing facts should trigger clarification, escalation, or restraint.

09 | SOURCE & FRAME TRANSPARENCY

Expose not only citations, but how authorities are selected and presented.

10 | CONTESTABILITY

Provide meaningful routes to challenge, correct, review, and escalate.

TRUSTWORTHINESS IS NOT ONE MODEL PROPERTY. It is an architecture of technical and institutional controls.

Accountability. Responsibility. Transparency.

The question is no longer only:
'Can AI reason like a lawyer?'

**What must remain
VISIBLE,
ATTRIBUTABLE,
and
CONTESTABLE
when AI sounds legally authoritative?**