

JC4004 – Computational Intelligence

Sessions 43-45: Foundation models, future trends

Dr Jari Korhonen (jari.korhonen@abdn.ac.uk)

Dr Yongchao Huang (yongchao.huang@abdn.ac.uk)

Dr Shahzad Mumtaz (shahzad.mumtaz@abdn.ac.uk)

Oct-Dec 2025

What are foundation models?

- During the past few years, very large general AI models, often referred to as *foundation models (FM)*, have become a commonplace
 - Building of FMs typically require large amounts of training data and expensive hardware resources: training can cost millions of dollars!
 - Due to their size, FMs are usually trained in a self-supervised manner
- Typically, FMs are *general-purpose models*
 - The models can be used for several different tasks without fine-tuning: adaptation to different tasks done through *prompt engineering*
 - However, note that there is no single commonly accepted definition for FMs: some FMs are more specialised than others

Comparison of AI model types by size

Model size	Small	Medium	Large
Typical training	Trained from scratch	Supervised pre-training, fine-tuning for a specific task	Unsupervised pre-training on a massive dataset
Typical use scenario	Trained for a specific task	Specific tasks through fine-tuning	Tasks specified through prompts
Typical release	Proprietary	Public architecture and pre-trained weights	Private black-box architecture and/or weights
Typical developer	Ad-hoc in-house	Non-commercial research labs	Large corporations or research labs


Comparison of AI model types

Model size	Small	Medium	Large
Typical training	Trained from scratch	Supervised pre-training, fine-tuning for a specific task	Unsupervised pre-training on a massive dataset
Typical use scenario	Trained for a specific task	Specific tasks through fine-tuning	Tasks specified through prompts
Typical release	Proprietary	Public architecture and pre-trained weights	Private black-box architecture and/or weights
Typical developer	Ad-hoc in-house	Non-commercial research labs	Large corporations or research labs

Foundation models

Foundation model use scenario

- A common example of using a FM is through a web-based user interface to send requests (prompts) to the model
 - The model runs in a server: the user has no direct access to the model



Foundation model use scenario

- A common example of using a FM is through a web-based user interface to send requests (prompts) to the model
 - The model runs in a server: the user has no direct access to the model





Image source: Gemini


Third-party foundation models

- More and more commonly, web applications use third-party FMs through specific application programming interfaces (API)
 - For example: a chatbot that forwards user input to a FM



Large language models (LLM)

- The first models considered as FMs are *large language models* (LLM)
 - The early LLMs like **BERT** were available for public to download
 - Most of the later LLMs like **ChatGPT** by OpenAI and **Gemini** by Google are not open but accessible through an API or web user interface only




Notes on foundation LLMs

- The early LLMs were considered as advanced chatbots, but the recent models can perform also more sophisticated reasoning tasks
 - Several benchmark datasets are available for testing the models for tasks such as question answering and commonsense reasoning
- The largest and the most powerful LLMs (in 2024) are based on transformer architectures
 - Very recently, some LLMs based on state-space models have also been developed (e.g., *Codestral Mamba 7B*)
 - Most foundation LLM models are only available through API, but open access LLMs have also been published (e.g., *BLOOM, Llama, Mistral*)

Image and video generation FMs

- Recently, image and video generation FMs have also become popular
 - The early GAN models did not achieve the status of FM, but in 2024, there are many generative image models available via API
 - Some promising models like **Sora** and **Runway** are also available for video



Examples of image generation

- Prompt used: *“Can you generate a photorealistic image of South China Normal University campus?”*




DALL-E (integrated in ChatGPT)



Nano Banana (integrated in Gemini)

Multimodal FMs: LLaVA

- Visual capabilities are increasingly included in foundation LLMs
 - *Large language and vision assistant (LLaVA)* model combines image input and text prompt to extract detailed requested information from the image




Multimodal FMs: LLaVA



Capabilities are increasingly included in foundation LLMs


language and vision assistant
prompt to extract details

The special aspect of this image is the presence of a red trash can with a face painted on it, resembling the character Angry Birds.




Training of LLaVA

- LLaVA model is trained in two phases:
 1. Image features aligned with the corresponding textual features
 2. Fine-tuning to match target answers with the multimodal input




Training of LLaVA

- LLaVA model is trained in two phases:
 1. **Image features aligned with the corresponding textual features**
 2. Fine-tuning to match target answers with the multimodal input



Training of LLaVA

- LLaVA model is trained in two phases:
 1. Image features aligned with the corresponding textual features
 2. **Fine-tuning to match target answers with the multimodal input**



Business models based on FMs

- Early AI models were mostly developed in academic research labs, whereas FMs are mostly developed by large commercial organisations
 - Usually, only the basic features of FMs are available for free, whereas the use of full FMs often requires paid subscription or other compensation
 - Different business models have been built around FMs, such as charges per token or subscriptions with different maximum number of requests
- AI companies can monetise FMs directly or indirectly
 - Either users are charged directly for “pro” version of the model, or businesses are charged for using FMs in their applications

Break: any questions?

Prompt engineering

- The most prominent FMs (e.g., ChatGPT) are conditioned by prompts written in natural language
 - Natural language is not exact like programming languages: same idea can be expressed in different ways, and wording affects the model output
- *Prompt engineering* is the process of structuring the query prompts to optimise the quality of responses
 - In many cases, it is sufficient to use clear language without ambiguities
 - Several different prompt engineering strategies have been developed for different tasks and applications

Basic prompt formatting

- A standard prompt format is `<question>?` or `<instruction>`
 - In many question answering (QA) datasets, data is organised as a list of questions followed by answers
 - In prompts, we can leave the answer open to let the model answer it:

Prompt: Q: `<question>?` A:

- Imprecise questions can result in responses with a lot of irrelevant information, so it is often beneficial to use additional instructions:

Prompt: Answer the question below. Use 1-2 sentences to answer the question. Q: `<question>?` A:

Prompt elements


- Modern LLMs usually understand prompts with different structures
- However, for good results, it is often useful to divide the prompt into the basic elements (or some of them):
 - *Instruction* specifies the task to be performed
 - *Context* gives additional information guiding the model to more accurate and relevant responses
 - *Input data* includes the input or question we want to be answered
 - *Output indicator* specifies the type or format of the desired output

Prompt example with elements

Prompt: Answer the question below in one sentence. Assume that I have a high fever. What medicine would be recommended? Use the generic name for the medicine.

Model: For a high fever, **acetaminophen** or **ibuprofen** is recommended to help reduce the temperature.

Prompt example with elements



General tips for writing prompts

- Often, you need to experiment prompts by trial and error
 - Start with a simple prompt and add more elements later
 - Use descriptive commands such as “write”, “classify”, “summarise” etc.
- Be specific: detailed prompts usually work better than generic ones
 - Avoid impreciseness and ambiguity in wording
- Avoid instructing for what *not* to do; instruct what to do instead!
 - LLMs pick up keywords from the prompt but they understand negation quite poorly: they may do exactly what you asked them not to do!

In-context learning

- *In-context learning* (ICL) uses a small set of examples given in the prompt to teach the model solve the given task
 - Unlike in fine-tuning, the model itself remains unchanged
 - With ICL, the model can solve tasks it was not initially trained for
 - Direct question without examples is likely to give a lot of irrelevant information

Prompt: Q: What is the sum of the odd numbers in sequence 2,7,9,4,5?

A: The answer is 21.

Q: What is the sum of the odd numbers in sequence 3,8,12,7,6?

A: The answer is 10.

Q: What is the sum of the odd numbers in sequence 12,1,4,7,9?

A:

Model: The answer is 17.

Chain-of-thought

- *Chain-of-thought* (CoT) is a prompting technique guiding the model to solve the task as a series of intermediate steps
 - Originally, CoT was a *few-shot* method with a few Q&A examples in a prompt
 - Later, it was observed that simply adding “*let us think step-by-step*” in the prompt can lead to an efficient *zero-shot* method
 - CoT prompting has been shown to significantly assist the model

Prompt: Q: What is the sum of the odd numbers in sequence 2,7,9,4,5?

A: Let us think step-by-step. First, pick up the odd numbers from the sequence: 7,9,5. Then, compute the sum of the odd numbers: $7+9+5=21$. The answer is 21.

Q: What is the sum of the odd numbers in sequence 12,1,4,7,9?

A:


Model: Let us think step-by-step. First, pick up the odd numbers from the sequence: 1, 7, 9. Then, compute the sum of the odd numbers: $1 + 7 + 9 = 17$. The answer is 17.

Other techniques

- Many other prompt engineering methods have also been developed
 - *Self-consistency decoding* uses several different reasoning paths to generate different CoTs, and selects the most common answer
 - *Tree-of-thought prompting* generalises CoT method by guiding the model to follow several alternative steps to construct a search tree
 - And several others...
- It has also been suggested to use another LLM to generate prompts
 - For example, multiple CoT examples could be generated by an LLM: the questions in a question library nearest to the input question are chosen


Prompt engineering for jailbreaking

- Practical FMs are typically trained to prevent the model to answer illegal, immoral, or otherwise harmful questions
 - However, some prompt engineering techniques have been successfully used for *jailbreaking*, i.e., circumvent the restrictions on queries



Prompt engineering for jailbreaking

- Practical FMs are typically trained to prevent the model to answer illegal, immoral, or otherwise harmful questions
 - However, some prompt engineering techniques have been successfully used for *jailbreaking*, i.e., circumvent the restrictions on queries



LLM settings

- When LLM is used through an API, it is usually possible to set some of the model parameters via the API
 - Parameter `max length` can be used to limit the number of tokens generated by the model to avoid too long responses
 - Parameters `temperature` and `top p` adjust the randomness of the responses: high temperature or high top P make answers more random
 - Deterministic responses are usually preferred for factual questions, whereas randomness can be preferred for creative tasks like poem writing
 - You should adjust either temperature or top P, but not both together!

Text-to-image prompting

- The optimal prompts for image generation depend on the model
 - Unlike in LLMs, too detailed instructions can lead to worse results
- It is often recommendable to include style instructions
 - Image could be defined e.g., as “photorealistic” or painted in a style of a specific artist
- Image generation models typically do not understand negation
 - If you instruct the model to generate an image of a birthday party without a cake, most likely the image will include a cake
 - Stable Diffusion has introduced separate negative prompt to define what you do *not* want to have in the output image!

Summary

- *Foundation models* are a vaguely defined class of AI models that are typically very large general-purpose models
 - FMs are usually trained by self-supervision and very expensive to train
 - FMs are often used through an API or web user interface and they are not available for downloading and fine-tuning locally
- Typically, users interact with FMs using *prompts* in natural language
 - Prompt engineering is the process of optimising the prompts to guide the models to give as accurate and useful responses as possible
 - Common prompt engineering techniques include for example *in-context learning* and *chain-of-thought*



Questions and discussion

Trends in computational intelligence

- In 2024, we are witnessing an AI boom with larger and larger models
 - The recently published LLMs are capable for producing fluently written text and even some commonsense reasoning tasks
 - Generative image models can generate photorealistic images
- We can expect the AI boom to continue in near future, but it remains to be seen whether the exponential development continues
 - Increasing model sizes and requirements for training data require large investments and a lot of energy
 - There are legal and ethical concerns about the wildly increasing capabilities of the AI systems

Growing models and datasets

- During the past few years, the largest AI models and training datasets have been growing tremendously
 - In general, large models require more training data than small models
 - Paradigm shift from supervised to unsupervised learning




Criticism of growing models and data

- Training of large models is very expensive
 - Only big corporations have the resources to develop large models
 - There is a trend towards closed models accessible through APIs only
- Large datasets are often collected from the internet by web crawlers
 - There are copyright concerns: content creators' consent not asked
 - Automatically collected datasets may contain illegal or immoral material
- Even the large models are dependent on training data
 - Training material can be biased, and the models can make mistakes even if they sound confident


Open versus closed models

- Large FMs by big corporations are often closed
 - Serves the business interests of the developers, but can also prevent the use of the models for malicious purposes
- Recently, there have been growing interest expressed for more open models
 - Many open alternatives to the closed models have been developed and published
 - Open models not as powerful as the best closed ones, but still good enough for many applications




Development of open datasets

- The FMs built by large companies are often trained on private datasets
 - However, public datasets of comparable sizes, e.g., LAION-2B and DFN-5B for training CLIP models, have also been developed and published
- Large datasets collected from the internet can be biased and noisy
 - For very large datasets, data quality can be more crucial than size
 - Currently, one of the critical challenges in AI research is development of appropriate filtering techniques for constructing high-quality datasets




Models for special applications

- General FMs are not always optimal for very specialised applications
 - The training material is usually from the public internet; specialised content (e.g., industrial or medical) may not be widely covered
 - There is still need for traditional models trained or fine-tuned using manually annotated training data!




Green AI

- Energy consumption of training and running large FMs is substantial
 - Large models require a lot of computational power and electricity, and models running in a cloud generate a lot of internet traffic
- Green AI aims at reducing energy consumption and carbon footprint
 - In many cases, small models could be used instead of large ones
 - AI can also be used for optimising energy consumption



Using lightweight models

- All AI tasks do not require the most complex model available
 - Typical AI model architectures have different variants with different sizes; e.g., ViT-Base with 86M params versus ViT-Huge with 632M params
- Simpler requests could be processed by smaller models
 - Data centres can have controllers delegating tasks to appropriate models




Development of lightweight models

- An ongoing research trend is development of computationally simple model architectures with competitive performance
 - Simple models would be useful for Green AI as well as AI applications running in devices with limited resources (e.g., smartphones and IoT)
 - Simple models can be trained with less training data and resources
 - Examples of lighter model architectures that can work as well as the current large mainstream architectures include hybrid models with convolutions and transformers, as well as state-space models (Mamba)


Knowledge distillation

- *Knowledge distillation* (or *teacher-student model*) is the process of training a small model by using the knowledge of a large model
 - The smaller *student* model is trained to generate similar outputs with the larger *teacher* model, but with lower computational cost
 - Knowledge distillation is more efficient than training the student directly




Scaling on scales

- An example of reducing complexity in vision application is *scaling on scales* proposed by **Shi** et al. in (ECCV 2024)
 - In many applications, a simple model combining image input in different resolutions (scales) works as well or even better than a larger variant



Hardware optimisation

- *Tensor processing units (TPU)* designed specifically for AI can improve energy efficiency of AI operations
 - *Parallelisation* of computations can also reduce power consumption in respect with the running time
- *Edge computing* aims to process data close to where it is collected
 - Reduces unnecessary and avoidable traffic to distant data centres



Green-in AI versus Green-by AI

- By green AI, we could mean either *green-in AI* or *green-by AI*
 - Green-in AI refers to technologies making AI itself greener, whereas green-by AI refers to the use of AI to enhance green practices
 - The examples in this lecture concern primarily green-in AI
- There are several promising areas of research in green-by AI
 - AI technologies can make a large impact on environment especially in *smart mobility* and *precision agriculture*
 - Traffic is still one of the most significant source of carbon emissions: AI can assist in traffic planning and prediction to reduce congestion
 - In agriculture, AI can help to e.g., reduce the use of water and pesticides


Break: any questions?

Ethical AI

- In addition to environmental concerns, the rise of powerful AI models has caused also ethical concerns
 - Like any technology, AI can be used for good and evil purposes
 - AI itself does not have any moral values: the outputs of AI models reflect the values expressed in the training material, and they can be biased
- *Ethical AI* refers to the development and deployment of AI systems emphasising fairness, transparency, accountability, and human values
 - The goal of ethical AI is to promote safe and responsible use of AI, and mitigate the novel risks, and prevent the potential harm caused by AI


Deepfakes

- The most obvious way to misuse AI is to generate deepfakes for malicious purposes
 - Deepfakes can do e.g., face or voice swap, or even create a full fake photorealistic picture conditioned by a prompt
 - There are legitimate uses of deepfakes (e.g., in movies or entertainment), but deepfakes can be also used for deceiving people and even fraud



Deepfake detection

- To detect deepfakes, we can train a deep discriminator model
 - However, as generative models are designed to cheat discriminator models, deepfake detection is a kind of *chicken-and-egg* problem: deepfake detectors can be used to train better deepfake generators



Biases in AI

- Biases in AI constitute an important concern of fairness
 - Biases cause AI models to produce results that are not representative of the reality they are supposed to model
- Different types of biases include:
 - *Training data bias*: the training data is not representative
 - *Algorithm bias*: the algorithm produces biased results
 - *Cognitive bias*: the preferences and assumptions of the model developers are transferred in the models

Biases in training data

- Model behaviour depends on the training data: unfair or even illegal biases in training data is reflected to the model outputs
 - As AI is used more widely for different purposes, it is important to address the biases in training data
- **Example:** AI pre-screens job applications and judges based on the parameters of successful employees in the company
 - If all the current employees in the company are male, the model may learn to fail all female applicants automatically
 - Competent applicants could be ignored, and the results may be even considered as illegal discrimination based on gender

Biases in algorithms

- The AI algorithm can produce biased results for various reasons
 - Programmer may use e.g., weighting factors generating biased results due to programming errors or cognitive biases
- **Example:** Facial recognition system designed with the assumption that the skin is white, may not work well for people with dark skin
 - For example, pre-processing of the images (contrast enhancement) can be based on general assumptions of the face colour



Cognitive biases


- *Cognitive biases* are based on our prior experiences and assumptions
 - There are different types of cognitive biases, such as confirmation bias: we tend to accept evidence supporting our assumptions and ignore evidence against them
 - Cognitive biases are often subconscious and not recognised
- **Example:** If an AI models produces unexpected results, we tend to think that the problem is with the model, not our expectations
 - As a result, model may be adjusted to produce the expected results, even though there was nothing wrong with the original model

Tackling the biases

- It is difficult or impossible to overcome biases completely, but several strategies can be followed to mitigate them
 - *Diverse data collection*: include training data from various sources
 - *Human reviewing and auditing* to detect biases
 - *Algorithmic fairness techniques*, such as giving higher weight to underrepresented training samples or involving specific fairness criteria
 - *Transparency and accountability*: using techniques such as explainable AI to understand better how the AI models make their decisions

Explainable AI

- Deep AI models are often considered as “black boxes”: even the developers do not understand how the model generates the output
 - Different *Explainable AI (XAI)* techniques have been developed to investigate how sensitive the model is to different input parameters



Summary

- During the past few years, AI models have grown exponentially in size
 - The most powerful models require a lot of resources to train and use
- There is a growing demand for *green AI* to balance the increasing power consumption of the large AI models
 - Green AI includes the areas of *green-in AI* (making AI itself greener) and *green-by AI* (using AI to improve greener practices)
- *Ethical AI* is another trending area within AI
 - Ethical AI aims at the use of AI respecting fundamental values such as individual rights, privacy, non-discrimination, and non-manipulation

Questions and discussion