



**JAX: Functional
Transformations for
High-Performance
ML (2018) -
developed by
Google**

- Transform numerical Python/NumPy functions
- Functional programming paradigm; immutable arrays
- XLA (accelerated linear algebra)-backed just-in-time compilation for CPUs/GPUs/TPUs

Why It Matters:

- Research velocity with concise, pure functions
- Performance: kernel fusion, layout optimization, device placement via XLA
- Scales from single device to multi-accelerator setups (TPUs/GPUs)

Just-in-Time (JIT) compilation is a method for executing computer code that translates it into native machine code at runtime. It is a hybrid approach that combines the best aspects of traditional interpreters (flexibility) and static, ahead-of-time (AOT) compilers (performance).

XLA, which stands for Accelerated Linear Algebra, is an open-source compiler infrastructure designed to improve the performance of machine learning models and other numerical computations, particularly those involving large-scale array operations.

Framework Comparison Table

	TensorFlow	PyTorch	JAX
Graph Type	Static (1.x), Dynamic (2.x)	Dynamic	Functional transformations
Programming Model	Imperative (2.x), Symbolic (1.x)	Imperative	Functional
Core Data Structure	Tensor (mutable)	Tensor (mutable)	Array (immutable)
Execution Mode	Eager (2.x default), Graph	Eager	Just-in-time compilation
Automatic Differentiation	Reverse mode	Reverse mode	Forward and Reverse mode
Hardware Acceleration	CPU, GPU, TPU	CPU, GPU	CPU, GPU, TPU

Eager execution in TensorFlow is an imperative programming environment where operations are evaluated immediately as they are called, rather than building a computational graph to be run later. This contrasts with TensorFlow's traditional graph mode, where operations are first added to a symbolic graph and then executed in a TensorFlow session.

Quiz

1) Why do deep learning workloads map well to GPUs?

- A. GPUs have larger caches than CPUs
- B. GPUs excel at massively parallel dense linear algebra like GEMMs/convolutions
- C. GPUs run higher clock speeds than CPUs
- D. GPUs avoid using memory entirely

2) In data-parallel training, what communication step typically synchronizes model replicas each iteration?

- A. Parameter server RPCs
- B. All-to-all exchange of activations
- C. All-reduce of gradients
- D. Broadcasting the dataset

3) Which statement best describes JAX's programming model?

- A. Imperative mutable tensors with static graphs
- B. Pure functions over immutable arrays transformed by jit/grad/vmap/pjit
- C. Graphs defined in a separate configuration language
- D. Only forward-mode differentiation is supported

4) What benefit does kernel fusion provide on accelerators?

- A. Increases model accuracy
- B. Reduces op launch overhead and memory traffic by combining ops
- C. Eliminates the need for backpropagation
- D. Doubles device memory capacity

5) When scaling batch size across more GPUs, which training adjustment is commonly recommended?

- A. Decrease learning rate proportionally to batch size
- B. Keep learning rate fixed regardless of batch
- C. Use the linear scaling rule and warmup the learning rate
- D. Disable mixed precision to stabilize training

SESSION 2: REVERSE MODE AUTOMATIC DIFFERENTIATION

The Backbone of Neural Network Training

Key Insight: One scalar output (loss),
millions of parameters

Two-Phase Process:

- Forward pass: Compute and store intermediates
- Backward pass: Propagate derivatives from output to inputs

Perfect Match: Many inputs, few outputs

