




TensorFlow Lite Micro: Architecture Decisions

* During Interpretation, the runtime reads and executes a model **directly from a serialized format** (FlatBuffers) without generating new source code or recompiling the project for every model. [Read more on Flatbuffers here.](#)

Interpretation vs Code Generation

Choice: Model interpretation rather than code generation

Benefits of Interpretation:

- **Upgradability** - easy framework updates
- **Multiple models** support without code duplication
- **Model replacement** without recompilation
- **Inline data** from FlatBuffers format *

Project Generation Approach:

- **Copy required source files** without modification
- **IDE-specific project files** for easy building
- **External dependencies** included in project
- **Upgradable** through standard merge tools

Model File Format: FlatBuffers

FlatBuffers is a binary serialization format plus a schema compiler. You define a schema (what fields exist), then serialize data into a flat binary layout that can be read directly without parsing or allocating. Readers access fields via pointers/offsets—no unpacking step.

Why FlatBuffers?

- Zero-copy deserialization - direct memory access
- Same in-memory and serialized format
- Embedded directly in flash memory
- No parsing overhead

Schema Structure (schema.fbs):

```
table Model {  
  version: uint;  
  operator_codes: [OperatorCode];  
  subgraphs: [SubGraph];  
  buffers: [Buffer];  
}  
  
table SubGraph {  
  tensors: [Tensor];  
  operators: [Operator];  
  inputs: [int];  
  outputs: [int];  
}  
  
root_type Model;
```

Model Loading Example

Loading a Model:

```
// Map model data (no copying/parsing)
const tflite::Model* model =
    tflite::GetModel(g_tiny_conv_micro_features_model_data);

// Version check
if (model->version() != TFLITE_SCHEMA_VERSION) {
    error_reporter->Report("Schema version mismatch");
    return 1;
}

// Access subgraph
auto* subgraphs = model->subgraphs();
if (subgraphs->size() != 1) {
    error_reporter->Report("Only 1 subgraph supported");
    return 1;
}
subgraph_ = (*subgraphs)[0];
```

Key Point: Direct memory access, no data structure construction

Tensors and Buffers in FlatBuffers

Buffer: Raw data storage

```
table Buffer {  
  data: [ubyte] (force_align: 16); // 16-byte aligned raw data  
}
```

Tensor: Metadata + buffer reference

```
table Tensor {  
  shape: [int]; // Dimensions [batch, height,  
width, channels]  
  type: TensorType; // float32, int8, etc.  
  buffer: uint; // Index into buffers array  
  name: string; // For debugging  
  quantization: QuantizationParameters; // For quantized  
models  
}
```

Separation: Data (buffers) vs Metadata (tensors)